



Intel® Deep Learning Studio

Release Notes for Alpha Release

Intel Confidential—Do Not Reproduce
Revision 1.0, September 2018



Copyright © 2018 Intel Corporation. All rights reserved.

Intel and the Intel logo are trademarks of Intel Corporation in the U.S. and other countries.

* Other names and brands may be claimed as the property of others.

This document contains information on products and/or processes in development.

INFORMATION IN THIS DOCUMENT IS PROVIDED IN CONNECTION WITH INTEL PRODUCTS AND/OR SERVICES. NO LICENSE, EXPRESS OR IMPLIED, BY ESTOPPEL OR OTHERWISE, TO ANY INTELLECTUAL PROPERTY RIGHTS IS GRANTED BY THIS DOCUMENT. EXCEPT AS PROVIDED IN AN EXECUTED AGREEMENT, INTEL ASSUMES NO LIABILITY WHATSOEVER AND INTEL DISCLAIMS ANY EXPRESS OR IMPLIED WARRANTY, RELATING TO SALE AND/OR USE OF INTEL PRODUCTS AND SERVICES INCLUDING LIABILITY OR WARRANTIES RELATING TO FITNESS FOR A PARTICULAR PURPOSE, MERCHANTABILITY, OR INFRINGEMENT OF ANY PATENT, COPYRIGHT, OR OTHER INTELLECTUAL PROPERTY RIGHT.

UNLESS OTHERWISE AGREED IN A SIGNED WRITING BY INTEL, THE INTEL PRODUCTS AND SERVICES ARE NOT DESIGNED NOR INTENDED FOR ANY APPLICATION IN WHICH THE FAILURE OF THE INTEL PRODUCT COULD CREATE A SITUATION WHERE PERSONAL INJURY OR DEATH MAY OCCUR.

Intel may make changes to specifications and product /process descriptions at any time, without notice. Designers must not rely on the absence or characteristics of any features or instructions marked "reserved" or "undefined." Intel reserves these for future definition and shall have no responsibility whatsoever for conflicts or incompatibilities arising from future changes to them. The information here is subject to change without notice. Do not finalize a design with this information.

The products and services described in this document may contain defects or errors known as errata which may cause deviations from published specifications. Current characterized errata are available on request.

If the end user enables remote access for Intel, Intel may be able to access system log files, which includes information such as IP address, User ID and Tenant ID, which may be considered personal information under the laws of certain jurisdictions. This information will be used only as necessary to diagnose, remediate and fix reported issues.



Document Revision History

Revision Number	Date	Comments (Latest on Top)
1.0	September 2018	Initial Release Alpha.



Contents

Document Revision History.....	3
Product Summary	5
Intel® DL Studio CLI Supported Operating Systems.....	5
Intel DL Studio Platform Installer Supported Operating Systems	5
Supported Browsers	5
Supported Hardware	5
Included Software Components.....	5
Known Issues.....	7



Product Summary

The Intel® Deep Learning Studio software application lets you run deep learning model training experiments and view the results of those experiments. You can run jobs on Kubernetes* and view results on TensorBoard*. You can use existing data sets, use your own data, or downloaded data from online sources, and can use existing templates or customize templates. You can run multiple experiments on multiple nodes of a Kubernetes cluster.

Intel® DL Studio CLI Supported Operating Systems

- macOS* High Sierra
- Ubuntu* 16.04
- Windows* 10

Intel DL Studio Platform Installer Supported Operating Systems

- Red Hat* Enterprise Linux* 7.5
- Ubuntu 16.04
- CentOS* 7.5

Supported Browsers

- Google* Chrome* 62+ for Windows, macOS* and Linux
- Additional browsers, such as Mozilla* Firefox* and Apple* Safari* may work, however, they have not been validated.

Supported Hardware

Optimized for up to 16 Intel® Xeon® processors (formerly Intel® microarchitecture code name Skylake).

Included Software Components

This release of the Intel DL Studio and platform software, version: 1.0.0-20180920074938, includes the following software components.

Name	Version	Project Link
kube-proxy	1.10.6	http://gcr.io/google-containers/kube-proxy-amd64
pause	3.1	http://gcr.io/google-containers/pause-amd64
tiller	2.9.1	https://github.com/helm/helm
addon-resizer	1.7	http://k8s.gcr.io/addon-resizer
heapster	1.4.2	https://github.com/kubernetes/heapster



Name	Version	Project Link
defaultbackend	1.4	https://github.com/kubernetes/ingress-nginx
ingress	0.14.0	http://quay.io/kubernetes-ingress-controller/nginx-ingress-controller
kube-dns	1.14.8	https://github.com/kubernetes/dns
dnsmasq-nanny	1.14.8	https://github.com/kubernetes/dns
dns-sidecar	1.14.8	https://github.com/kubernetes/dns
etcd	3.3.9	https://github.com/coreos/etcd
dashboard	1.8.3	https://k8s.gcr.io/kubernetes-dashboard-amd64
registry	2.6.2	https://github.com/docker/distribution
kubectrl	1.10.6	https://github.com/kubernetes/kubernetes/tree/master/pkg/kubectrl
elasticsearch	6.2.3	https://github.com/elastic/elasticsearch
fluentd	0.12	https://www.fluentd.org/
redsocks	0.5	https://github.com/darkk/redsocks
tensorflow	1.9.0	https://github.com/tensorflow/tensorflow
nginx	1.14.0	https://www.nginx.com/
mkl-dnn	0.14	https://github.com/intel/mkl-dnn



Known Issues

Defect ID	Details
CAN-1195	[dlctl] Other error during starting application – This appears sporadically while getting logs from the platform. Workaround: Retry getting logs from the platform.
CAN-1194	During multi-node training it was observed that one worker finished training 4-5 hours after other nodes. Workaround: N/A
CAN-1166	An experiment cancelled by pod status deletes all pods in the user's namespace. Workaround: N/A
CAN-1144	[dlctl] An experiment that has been cancelled already cannot be purged. Workaround: N/A
CAN-1104	Sometimes when Tensorflow is launched, it attaches to a non-existing log folder. Workaround: Retry launching TensorBoard*.
CAN-1085	Pod is started again after physical node reboot even though an experiment has been completed. Workaround: N/A
CAN-1078	All cores of the node are not being utilized by multi-node trainings, resulting in degraded performance on multi-node experiments. Workaround: N/A
CAN-1061	[dlctl] <code>dlctl experiment list</code> command does not show requested resources for queued experiments. Workaround: Use the <code>dlctl experiment view</code> command to show requested resources for an experiment.
CAN-1048	When an experiment node is stopped/turned off, the experiment is not started on another node. Workaround: N/A
CAN-1037	[dlctl] Status of inference instance should be FAILED and not QUEUED in case of <code>Init:CrashLoopBackOff</code> status of pod. Workaround: N/A
CAN-914	TensorBoard service occasionally throws 503 Service Unavailable. Workaround: Refresh the page. The session was created successfully, however the server was still initializing.



Defect ID	Details
CAN-906	[dlscctl] Creating interactive experiment with the same name as already canceled experiment causes error. Workaround: Create an interactive notebook again using a new name.
CAN-896	Experiment details cannot be displayed for experiments that failed during submit phase. Workaround: N/A
CAN-895	When submitting an experiment using Windows 10, a UnicodeEncodeError exception may be displayed. Workaround: N/A
CAN-891	[dlscctl] Failed to load TensorBoard in case of lack of 127.0.0.1 in no_proxy settings. Workaround: If the environment is configured with a proxy, then add 127.0.0.1 address to the no_proxy environment variable.
CAN-616	User with the same name cannot be created directly after its removal. Workaround: Create a username with new name.